


# Chapter 4: Threads & Concurrency

---





# SANITY CHECKING!!!!!!!

---

- Process versus thread?
- Race condition?
- Concurrency control?
- Concurrency versus parallelism?





# ADMIN...

---

- GREAT WORK!!!!
- A1 is AWESOME!
  - Still marking...
- MIDTERM looks GOOD!
  - Marking will take forever, but stay tuned!
- A2 Your first concurrency adventure is posted!
  - NOT using Pthreads YET!
  - It is a bit mind boggling!!!!
  - Understanding what can go wrong and why!!!!





# (BATTLE)WORM!

- In this part of the assignment, **you will implement** a simple system for running a series of **tasks** within one process.
  - Tasks in this system will run **separate functions**, and your scheduler will need to **switch between them**.
- These tasks will have the ability to perform **blocking** operations. When a running task blocks, your scheduler will need to stop running the current task and select a new one to run in the meantime.
- You will **not use preemption**; your scheduler only switches tasks when the currently-running task **blocks**. This is sometimes referred to as *cooperative scheduling*, since it relies on all the tasks to cooperate by blocking periodically.
- While the system you implement could be used for other purposes, you will be implementing it specifically to support the game **Worm!**, a clone of the classic game *Snake*.
- NOT AT ALL LIKE BATTLESNAKE... but a start!
- DOES NOT USE PTHREADS!!!!!!





# Details

---

- This game implementation uses several **tasks**:
  - A main task that starts up the game
  - A task to update the worm state in the game board
  - A task that redraws the board periodically
  - A task that reads input from the user and processes controls
  - A task that generates "apples" at random locations on the board
  - A task that updates existing "apples" by spinning them and removing them after some time
- Each task runs in a **loop** that contains a **blocking operation**.
- Your job is to implement a **scheduling system** that will run these tasks in **round-robin** fashion, switching between them as the currently-executing task blocks.
- The starter code has been provided, and take time to look it over as soon as you can!!!





# Assignment 2 Part 1 WORM!

---

- TRY to get this STARTED RIGHT AWAY!!!!
- TRY to get this working by next week...
- PART 2 (and 3) will be POSTED!
  
- YOU NEED TO THINK HARD!!!!!!
  - Tasks run until they BLOCK
  - When would each of the tasks
    - ▶ Run?
    - ▶ Block?
    - ▶ End?





# Questions & Answers

---

- How should `task_readchar` work?
  - If there's already a character to read, we could return it to the task immediately. However, if `getch()` returns ERR, **no input is available**.
  - First, you need to record **why the task is blocked** (it's waiting for input), and then invoke your scheduler to choose a new task to run.
  - Later, when the scheduler is choosing a task to run (at some point in the future) there will be an input character. The scheduler can then **`swapcontext()`** back to the task we blocked.





# Scheduler System Details

- You will be implementing what is known as *cooperative scheduling*.
  - This is a simple technique for implementing a scheduler without preemption. Once a task is started it will run until it issues a blocking operation or exits.
  - In our case, tasks will run until they *exit*, *wait* for another task, *sleep* for a fixed amount of time, or *wait* for user input. When you hit one of these points you should invoke the scheduler function to select and start another task to run.
- As part of this process, you will need to check to see if any previously-blocked tasks can now *unblock*.
  - This could happen because they are
    - ▶ waiting for tasks that have now completed,
    - ▶ their sleep timer has elapsed,
    - ▶ or user input is now available.







# Chapter 4: Threads

---

- Overview
- Multicore Programming
- Multithreading Models





# SANITY CHECKING!!!!!!!

---

- Process versus thread?
- Race condition?
- Concurrency control?
- Concurrency versus parallelism?





# Motivation


---

- Most modern applications are multithreaded
- Threads run within application
- Multiple tasks with the application can be implemented by separate threads
  - Update display
  - Fetch data
  - Spell checking
  - Answer a network request
- Process creation is heavy-weight while thread creation is light-weight
- Can simplify code, increase efficiency
- Kernels (NOW!) are generally multithreaded





# Multithreaded Server Architecture





# Benefits

---

- **Responsiveness** – may allow continued execution if part of process is blocked, especially important for user interfaces
- **Resource Sharing** – threads share resources of process, easier than shared memory or message passing
- **Economy** – cheaper than process creation, thread switching lower overhead than context switching
- **Scalability** – process can take advantage of multiprocessor architectures





# Multicore Programming

---

- **Multicore** or **multiprocessor** systems putting pressure on programmers, challenges include:
  - **Dividing activities**
  - **Balance**
  - **Data splitting**
  - **Data dependency**
  - **Testing and debugging**
- **Parallelism** implies a system can perform more than one task simultaneously
- **Concurrency** supports more than one task making progress
  - Single processor / core, scheduler providing concurrency





# Multicore Programming (Cont.)

---

## ■ Types of parallelism


- **Data parallelism** – distributes subsets of the same data across multiple cores, **same operation** on each
- **Task parallelism** – distributing threads across cores, each thread performing **unique** operation






# Concurrency vs. Parallelism

## ■ Concurrent execution on single-core system:



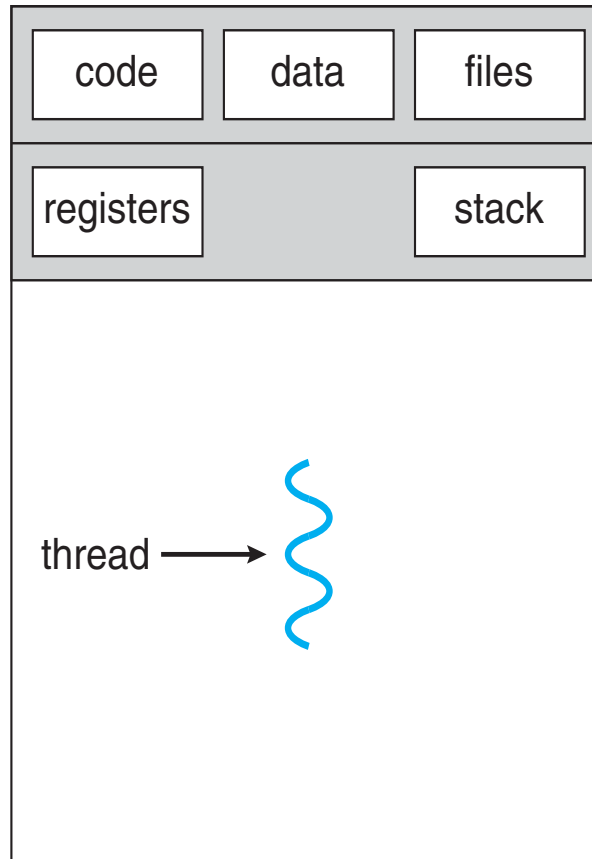
## ■ Parallelism on a multi-core system:








# Single and Multithreaded Processes



single-threaded process



multithreaded process





# User Threads and Kernel Threads

---

- **User threads** - management done by user-level threads library
- Three primary thread libraries:
  - POSIX **Pthreads**
  - Windows threads
  - Java threads
- **Kernel threads** - Supported by the Kernel
- Examples – virtually all general purpose operating systems, including:
  - Windows
  - Solaris
  - Linux
  - Tru64 UNIX
  - Mac OS X





# Multithreading Models

---

- Many-to-One
- One-to-One
- Many-to-Many





# Thread Libraries

---

- **Thread library** provides programmer with API for creating and managing threads
- Two primary ways of implementing
  - Library entirely in user space
  - Kernel-level library supported by the OS





# Pthreads

---

- May be provided either as user-level or kernel-level
- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- ***Specification***, not ***implementation***
- API specifies behavior of the thread library, implementation is up to development of the library
- Common in UNIX operating systems (Solaris, Linux, Mac OS X)





# BATTLEWORM!

---

- **Where in the scheduler should tasks actually be executed?**

Once your scheduler() function has selected a task to run, you call *swapcontext* to switch to that task. When the blocked task resumes, it will look like swapcontext() just returned zero.

- **What are task\_exit and the exit context for?**

The exit context tells the processor where to go when the task's main function returns. The starter code set this up so your tasks will call task\_exit when they finish. You should mark the task as DONE and then invoke the scheduler here.

- **How long should test3 take?**

It should be a total of ten seconds (plus a little random noise)





# BATTLEWORM!

---

## ■ When should the scheduler run?

When a task calls some function that blocks that task, invoke the scheduler. The scheduler should loop over tasks until it finds one that can run. Once you find the task, *swapcontext* to it. At this point, the scheduler is done.

## ■ Do you have any hints for `task_readchar`?

If a task blocks waiting for user input, the scheduler will only know if the task can resume if it calls `getch()`. If `getch()` returns `ERR` to the scheduler, the task can't run. If it returns something else, then you can run the task. Whatever value was returned by `getch()` should then be returned by `task_readchar()`.





# BATTLEWORM!

---

- **Is a context a list of instructions we need to run?**

Yes, in combination with the program source and memory. You can think of contexts as a specific spot we can return to in the middle of a set of steps.

- **A hint:**

If you are switching to a task that is the same as the current task, just return from the scheduler instead of calling `swapcontext()`.

