




CSC 360: Operating Systems

Tutorial #2



UNIX Shells and Process Creation

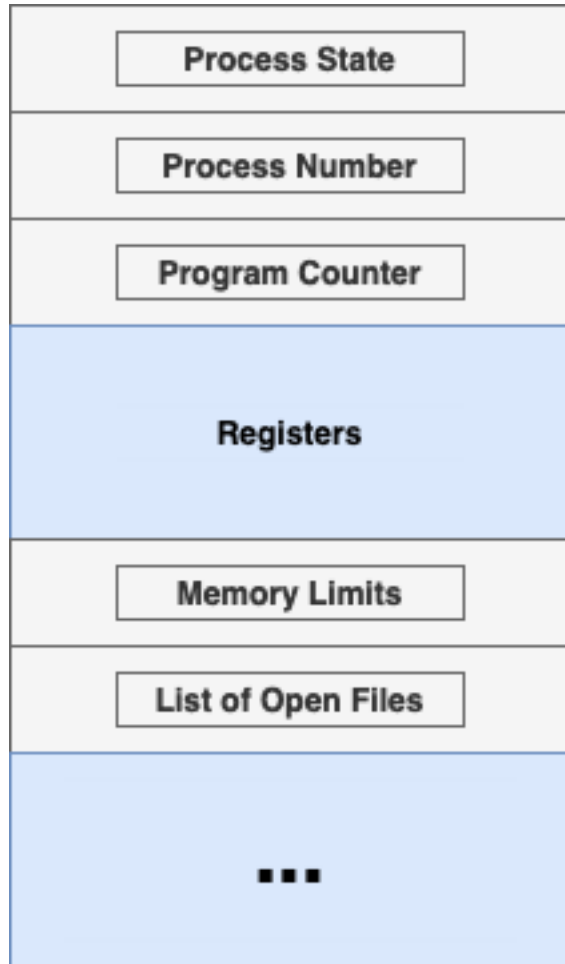
What is a Shell?

- ❑ Before there were graphical user interfaces, users had to write instructions to the computer on a keyboard.
- ❑ The earliest form of this was a teletype which was a typewriter with input and output. (We still have ASCII values 0x07 for bell and 0x0E for carriage return).
- ❑ A shell is a *command interpreter* for the operating system.
- ❑ It allows the user to execute programs as well as built in commands.
- ❑ When a user types a command, the shell parses it. It then looks in the \$PATH directories and built-in commands to determine what to do.
- ❑ When we run a program from the shell it spawns a new process and waits for it to finish.
- ❑ The shell is a user-level program that requests the OS to perform tasks on its behalf
- ❑ Examples: BASIC interpreters, Windows CMD.exe, Linux: sh, cshell, bash etc.

What is a Process?

- ❑ A process is an instance of a program that is currently executing on the system
- ❑ A program is simply a list of instructions that is stored permanently (on disk)
- ❑ When a process starts, these instructions are loaded into main memory (RAM). The operating system allocates each process its own address space.
- ❑ This ensures that processes cannot access the memory used by other processes
- ❑ Since many processes may be running at the same time, the OS must manage the sharing of system resources between processes
- ❑ The scheduler maintains state queues which determine whether a process is currently running on the CPU, waiting for the CPU, or doing something else. (just linked-lists etc.)
- ❑ Information about each process is stored in a *Process Control Block* (just a big struct)

What is in the PCB?



- ❑ The Linux kernel uses `task_struct` to implement the PCB
- ❑ It is simply a struct with various settings
- ❑ Some of these are:

```
pid_t pid;                // Used to identify the process
unsigned int cpu;          // The current CPU
int exit_code;             // Return value of process
struct task_struct __rcu *parent //Parent PCB
```

<https://github.com/torvalds/linux/blob/master/include/linux/sched.h>

How are processes created?

- ❑ Processes are started by *cloning* themselves from other processes
- ❑ A parent process will call the `fork()` system call.
- ❑ This causes a trap into kernel mode. The OS then runs appropriate code from a lookup table of system calls.
- ❑ The OS will then make an exact copy of the parent process's address space and PCB then adds the child process to the ready queue (starts execution).
- ❑ The only difference is that the child process has a different PID from its parent
- ❑ Both processes will resume executing from the instruction immediately following `fork()`
- ❑ In the parent, `fork()` returns the child's PID
- ❑ In the child, `fork()` returns 0
- ❑ This is how we differentiate between them

```
1  #include <sys/types.h>
2  #include <sys/wait.h>
3  #include <unistd.h>
4  #include <stdio.h>
5
6  int main()
7  {
8      pid_t pid;
9
10     pid = getpid();
11
12     printf("The PID of this process is: %d\n", pid);
13
14     pid = fork();
15     if (pid == 0) { //Returns PID of 0
16         // Child process
17         printf("This is the child process with fork returning PID: %d\n", pid);
18         sleep(10); //Just busy waits for 10 seconds
19     } else if (pid < 0) {
20         // Something went wrong
21         printf("Error forking:");
22     } else { //Returns PID of child
23         // Parent process
24         printf("This is the parent with fork returning PID: %d\n", pid);
25         wait(NULL); //Wait for child process to terminate before exiting
26     }
27 }
```

```
The PID of this process is: 12511
This is the parent with fork returning PID: 12513
This is the child process with fork returning PID: 0
```

How do we run a new program?

- ❑ We must use the `exec()` system call.
- ❑ After we have created a *cloned* process using `fork()`
- ❑ This copies the instructions (code) of a new program into the address space of the process that it is called from.


```
1  #include <sys/types.h>
2  #include <stdio.h>
3  #include <unistd.h>
4  #include <sys/wait.h>
5
6  int main()
7  {
8      pid_t pid;
9
10     /* fork a child process */
11     pid = fork();
12
13     if (pid < 0) { /* Something went wrong */
14         fprintf(stderr, "Fork failed\n");
15         return 1;
16     }
17     else if (pid == 0) { /* Child process */
18         execlp("/bin/ls", "ls", NULL);
19     }
20     else { /* Parent process */
21         /* Parent will wait for the child to complete */
22         wait(NULL);
23         printf("Child complete\n");
24     }
25
26     return 0;
27 }
28
```

What Should the Shell Do?

- ❑ Read an input string from the user
- ❑ Split the string into *tokens* and store in an array
- ❑ Analyze the token array to form commands, options and arguments
- ❑ Execute the command (built-in or look in path for program to run)
- ❑ Repeat...

Quiz

- ❑ Draw a tree diagram of the output of the following code:

```
<secret!>
```

In case you missed anything...

- ❑ Processes:

- ❑ http://man7.org/tlpi/download/TLPI-24-Process_Creation.pdf C Standard Library

- ❑ BASH:

- ❑ https://www.gnu.org/software/bash/manual/html_node/index.html#SEC_Contents

- ❑ Linux task_struct Process Control Block

- ❑ <https://github.com/torvalds/linux/blob/master/include/linux/sched.h>